

P8106-hw3

Renjie Wei

rw2844

3/19/2022

```
library(caret)
library(MASS)
library(mlbench)
library(pROC)
library(klaR)
library(glmnet)
library(pdp)
library(vip)
library(AppliedPredictiveModeling)
library(summarytools)
```

Split the dataset into two parts: training data (70%) and test data (30%).

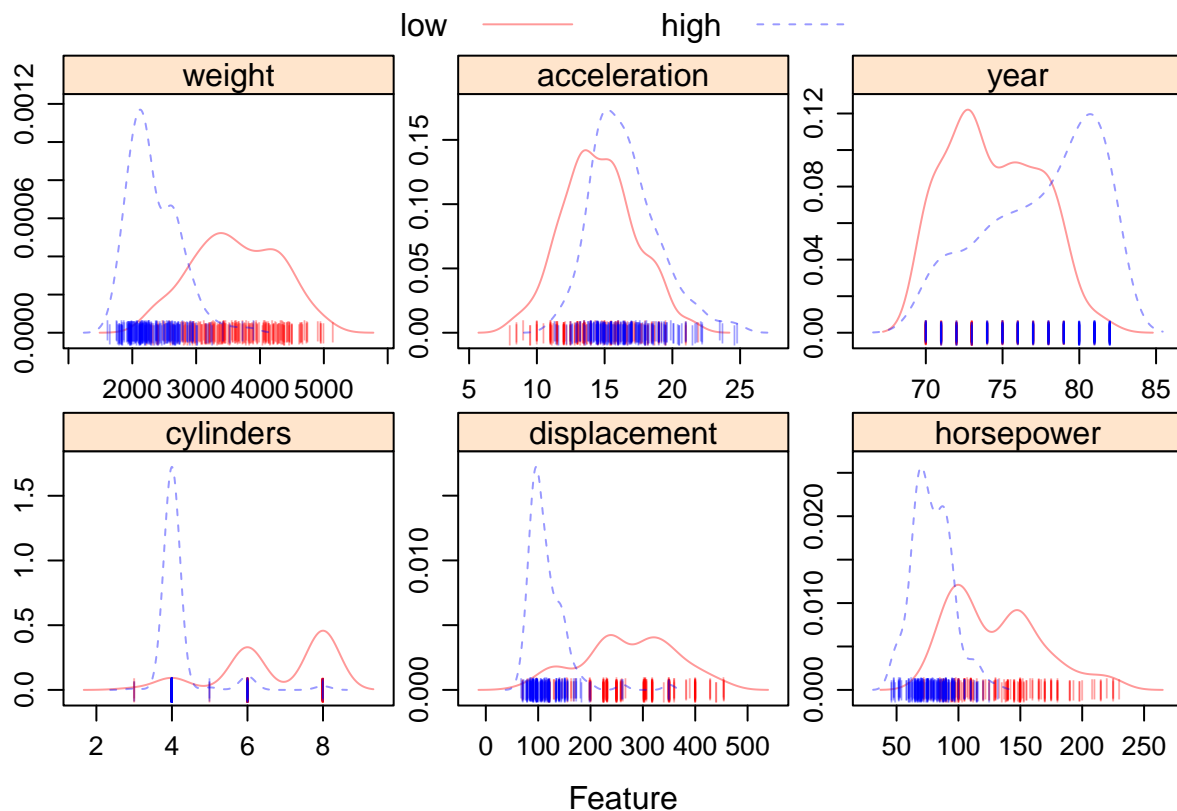
```
data <- read.csv("auto.csv")
data <- na.omit(data)
data$mpg_cat <- as.factor(data$mpg_cat)
data$mpg_cat <- relevel(data$mpg_cat, "low")
data$origin <- as.factor(data$origin)
set.seed(2022)
# sum(is.na(data)) = 0
rowTrain <- createDataPartition(y = data$mpg_cat, p = 0.7, list = FALSE)
```

(a) Produce some graphical or numerical summaries of the data

Answers

Here is a graphical summary for the continuous data

```
theme1 <- transparentTheme(trans = .4)
trellis.par.set(theme1)
featurePlot(x = data[, 1:6],
            y = data$mpg_cat,
            scales = list(x = list(relation = "free"),
                          y = list(relation = "free")),
            plot = "density", pch = "|",
            auto.key = list(columns = 2))
```



And here is a detailed summary for the total data.

```
dfSummary(data[, -1])
```

```
## Data Frame Summary
## Dimensions: 392 x 7
## Duplicates: 0
##
```

##	No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Valid
## 1	displacement	Mean (sd) : 194.4 (104.6)	81 distinct values	:	:	392
##	[numeric]	min < med < max:		:	:	(100.0%)
##		68 < 151 < 455		:	:	
##		IQR (CV) : 170.8 (0.5)		:	:	
##				:	:	
## 2	horsepower	Mean (sd) : 104.5 (38.5)	93 distinct values	:	:	392
##	[integer]	min < med < max:		:	:	(100.0%)
##		46 < 93.5 < 230		:	:	
##		IQR (CV) : 51 (0.4)		:	:	
##				:	:	
## 3	weight	Mean (sd) : 2977.6 (849.4)	346 distinct values	:	:	392
##	[integer]	min < med < max:		:	:	(100.0%)
##		1613 < 2803.5 < 5140		:	:	

```
##          IQR (CV) : 1389.5 (0.3)          : : : : :
##          : : : : : : :
##
## 4    acceleration    Mean (sd) : 15.5 (2.8)    95 distinct values    :    392
##      [numeric]      min < med < max:          : .    (100.0%)
##                      8 < 15.5 < 24.8          : : :
##                      IQR (CV) : 3.2 (0.2)      : : : :
##                      . : : : : : .
##
## 5    year            Mean (sd) : 76 (3.7)      13 distinct values    . : .    392
##      [integer]      min < med < max:          : : :    (100.0%)
##                      70 < 76 < 82            : : : : :
##                      IQR (CV) : 6 (0)         : : : : : : : :
##                      : : : : : : : :
##
## 6    origin          1. 1                    245 (62.5%)    I I I I I I I I I I    392
##      [factor]       2. 2                    68 (17.3%)    I I I    (100.0%)
##                      3. 3                    79 (20.2%)    I I I I
##
## 7    mpg_cat         1. low                   196 (50.0%)    I I I I I I I I I I    392
##      [factor]       2. high                  196 (50.0%)    I I I I I I I I I I    (100.0%)
## -----
```

- (b) Perform a logistic regression using the training data. Do any of the predictors appear to be statistically significant? If so, which ones? Compute the confusion matrix and overall fraction of correct predictions using the test data. Briefly explain what the confusion matrix is telling you.

Answers

```
ctrl <- trainControl(method = "repeatedcv", repeats = 5,
                      summaryFunction = twoClassSummary,
                      classProbs = TRUE)
set.seed(2022)
model.glm <- train(x = data[rowTrain,1:7],
                   y = data$mpg_cat[rowTrain],
                   method = "glm",
                   metric = "ROC",
                   trControl = ctrl)
summary(model.glm) # weight and year significant

##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.58449  -0.06036   0.00320   0.16299   2.80615
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -29.776125   8.020420  -3.713 0.000205 ***
## cylinders    0.159497   0.549610   0.290 0.771664
```

```
## displacement 0.014037 0.017660 0.795 0.426681
## horsepower -0.016364 0.029327 -0.558 0.576866
## weight -0.007198 0.001798 -4.003 6.24e-05 ***
## acceleration 0.114071 0.165285 0.690 0.490103
## year 0.605116 0.120409 5.025 5.02e-07 ***
## origin2 2.285179 0.978723 2.335 0.019551 *
## origin3 1.332239 0.927574 1.436 0.150928
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 382.617 on 275 degrees of freedom
## Residual deviance: 96.539 on 267 degrees of freedom
## AIC: 114.54
##
## Number of Fisher Scoring iterations: 8
```

From the model summary, it shows that `weight`, `year` and `origin2` (European) are statistically significant.

We set a cut-off value at 0.5 to build the confusion matrix

```
test.pred.prob <- predict(model.glm, newdata = data[-rowTrain,], type = "prob")[,2]
test.pred = rep("low", length(test.pred.prob))
test.pred[test.pred.prob>0.5] = "high"

confusionMatrix(data = relevel(as.factor(test.pred), "low"), reference = data$mpg_cat[-rowTrain], posit
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction low high
##      low  50   9
##      high   8  49
##
##           Accuracy : 0.8534
##           95% CI : (0.7758, 0.9122)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : 1.478e-15
##
##           Kappa : 0.7069
##
##      McNemar's Test P-Value : 1
##
##           Sensitivity : 0.8448
##           Specificity : 0.8621
##      Pos Pred Value : 0.8596
##      Neg Pred Value : 0.8475
##           Prevalence : 0.5000
##      Detection Rate : 0.4224
##      Detection Prevalence : 0.4914
##      Balanced Accuracy : 0.8534
##
##      'Positive' Class : high
```

##

```
#test.pred.class <- predict(model.glm, newdata = data[-rowTrain,], type = "raw")  
#confusionMatrix(data = test.pred.class, reference = data$mpg_cat[-rowTrain], positive = "high")
```


- The confusion matrix gives us the **Accuracy** : 0.8534, which means the misclassification rate is 1-Accuracy = 0.1466.
- **P-Value** [Acc > NIR] : 1.478e-15 means the accuracy is significantly larger than the no information rate which means our classifier is good.
- **Kappa** : 0.7069 evaluate the agreement between the predict result and observed result, and its quiet large, which means this agreement is not by chance.
- Both **Sensitivity** : 0.8448 and **Specificity** : 0.8621 are large, which also means our classifier is good.

```
glm.pred <- predict(model.glm, newdata = data[rowTrain,], type = "prob")[,2]  
roc.glm <- roc(data$mpg_cat[rowTrain], glm.pred)
```

(c) Train a multivariate adaptive regression spline (MARS) model using the training data.

Answers

```
set.seed(2022)  
model.mars <- train(x = data[rowTrain,1:7], y = data$mpg_cat[rowTrain], method = "earth", tuneGrid = exp  
plot(model.mars)
```



```
coef(model.mars$finalModel)
```

```
##              (Intercept)              h(250-displacement)
##              -7.980388e+00              7.287558e-02
##              h(year-72) h(4-cylinders) * h(250-displacement)
##              8.045225e-01              -1.166665e-01
## h(156-displacement) * h(year-72) h(250-displacement) * h(weight-2223)
##              -8.127365e-03              -4.147208e-05
```

```
mars.pred <- predict(model.mars, newdata = data[rowTrain,], type = "prob")[,2]
roc.mars <- roc(data$mpg_cat[rowTrain], mars.pred)
```

(d) Perform LDA using the training data. Plot the linear discriminants in LDA.



Answers

Fit the LDA model using MASS, and plot the linear discriminants.

```
set.seed(2022)
model.lda <- train(mpg_cat~.,
  data = data[rowTrain,],
  method = "lda",
  metric = "ROC",
  trControl = ctrl)
```

```
lda.pred <- predict(model.lda, newdata = data[rowTrain,], type = "prob" )[,2]
roc.lda <- roc(data$mpg_cat[rowTrain], lda.pred)

lda.fit <- lda(mpg_cat~., data = data, subset = rowTrain)
plot(lda.fit, col = as.numeric(data$mpg_cat[rowTrain]), abbrev = TRUE)
```




- (e) Which model will you use to predict the response variable? Plot its ROC curve using the test data. Report the AUC and the misclassification error rate.

Answers

To decide using which model to predict the response, I plot the ROC of the 3 models and compare their AUC on the train data.

```
auc <- c(roc.glm$auc[1], roc.mars$auc[1], roc.lda$auc[1])

modelNames <- c("glm","mars","lda")
ggroc(list(roc.glm, roc.mars, roc.lda), legacy.axes = TRUE) +
  scale_color_discrete(labels = paste0(modelNames, " (", round(auc,3),")"),
    name = "Models (AUC)" ) +
  geom_abline(intercept = 0, slope = 1, color = "grey")
```



The plot shows that `mars` model has the largest AUC compared to the rest. So I use `mars` to do the prediction. The model summary table below shows the same result.

```
res <- resamples(list(GLM = model.glm, MARS = model.mars, LDA = model.lda))
summary(res)
```


```
##
## Call:
## summary.resamples(object = res)
##
## Models: GLM, MARS, LDA
## Number of resamples: 50
##
## ROC
##           Min.   1st Qu.   Median     Mean   3rd Qu. Max. NA's
## GLM  0.8901099 0.9568289 0.9795918 0.9732055 0.9933281    1    0
## MARS 0.8622449 0.9649725 0.9843014 0.9785766 1.0000000    1    0
## LDA  0.8021978 0.9286771 0.9642857 0.9574749 0.9946942    1    0
##
## Sens
##           Min.   1st Qu.   Median     Mean   3rd Qu. Max. NA's
## GLM  0.7142857 0.8571429 0.9285714 0.9018681 0.9285714    1    0
## MARS 0.7857143 0.9230769 0.9285714 0.9263736 0.9285714    1    0
## LDA  0.7142857 0.7857143 0.8571429 0.8468132 0.9065934    1    0
##
## Spec
```



```
##           Min.    1st Qu.    Median    Mean 3rd Qu.  Max. NA's
## GLM  0.7692308 0.9230769 0.9285714 0.9372527      1      1      0
## MARS 0.8461538 0.9285714 0.9285714 0.9563736      1      1      0
## LDA  0.9230769 0.9285714 1.0000000 0.9723077      1      1      0
```

The ROC of mars model on the test set is shown below.

```
pred.test <- predict(model.mars, newdata = data[-rowTrain,], type = "prob")[,2]
roc.test <- roc(data$mpg_cat[-rowTrain], pred.test)
plot(roc.test, legacy.axes = TRUE, print.auc = TRUE)
plot(smooth(roc.test), col = 4, add = TRUE)
```



The test AUC is 0.948.

```
# set a cutoff Accuracy : 0.8793
pred.test.prob <- predict(model.mars, newdata = data[-rowTrain,], type = "prob")[,2]
pred.test.pred = rep("low", length(pred.test.prob))
pred.test.pred[pred.test.prob>0.5] = "high"
confusionMatrix(data = relevel(as.factor(pred.test.pred), "low"), reference = data$mpg_cat[-rowTrain], )

## Confusion Matrix and Statistics
##
##           Reference
## Prediction low high
##      low  51   7
```

```

##      high    7    51
##
##              Accuracy : 0.8793
##              95% CI : (0.8058, 0.9324)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.7586
##
##      McNemar's Test P-Value : 1
##
##              Sensitivity : 0.8793
##              Specificity : 0.8793
##              Pos Pred Value : 0.8793
##              Neg Pred Value : 0.8793
##              Prevalence : 0.5000
##              Detection Rate : 0.4397
##      Detection Prevalence : 0.5000
##              Balanced Accuracy : 0.8793
##
##      'Positive' Class : high
##
# use raw prediction Accuracy : 0.8793
#pred.test.class <- predict(model.mars, newdata = data[-rowTrain,], type = "raw")
#confusionMatrix(data = relevel(pred.test.class, "low"), reference = data$mpg_cat[-rowTrain], positive = "high")

```

The test Accuracy is 0.8793 so the misclassification error rate is $1 - \text{Accuracy} = 0.1207$.